



Green University of Bangladesh
Department of Computer Science and Engineering (CSE)
Faculty of Sciences and Engineering
Semester: (Fall, Year:2025), B.Sc. in CSE (Day)

Lab Report No: 1
Course Title: Data Mining Lab
Course Code: CSE-436 **Section:222-D1**

Lab Experiment Name: Feature Engineering techniques in Data Mining

Student Details

	Name	ID
1.	Md. Abdullah Al Moin	222902070

Submission Date : 13-10-2025
Course Teacher's Name : Md. Atikuzzaman

<u>Lab Report Status</u>	
Marks:	Signature:
Comments:	Date:

TITLE OF THE LAB REPORT EXPERIMENT:

Feature Engineering Techniques In Data Mining.

OBJECTIVES:

The main objective of this lab is to apply data preprocessing techniques such as imputation, outlier detection, feature engineering, feature scaling, and encoding. It also aims to build an automated workflow using pipelines to improve model performance and ensure efficient, consistent data handling during machine learning tasks.

PROCEDURE:

1. Load the dataset and inspect missing values.
2. Handle missing data using median imputation.
3. Encode categorical features with OneHotEncoder.
4. Split the dataset into training and testing sets.
5. Apply feature scaling for normalization.
6. Detect and remove outliers if needed.
7. Create a pipeline combining preprocessing and model training.

IMPLEMENTATION&TEST RESULT:

```
▶ from google.colab import drive
drive.mount('/content/drive')

→ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

s import pandas as pd

df = pd.read_csv("/content/drive/MyDrive/DataMining/oceanvoyageData.csv")
```

```
▶ !pip install pandas numpy scikit-learn matplotlib seaborn -q

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import sklearn

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, OneHotEncoder, PolynomialFeatures, KBinsDiscretizer, FunctionTransformer
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

import warnings
warnings.filterwarnings('ignore')

print(f"Pandas Version: {pd.__version__}")
print(f"NumPy Version: {np.__version__}")
print(f"Scikit-learn Version: {sklearn.__version__}")

→ Pandas Version: 2.2.2
NumPy Version: 2.0.2
Scikit-learn Version: 1.6.1
```

```
df.head()

→   Age      Fare TravelClass Gender EmbarkedPort SibSp Parch Survived
  0  28  231.376717          2    F       PortC     3     1      0
  1  40  18.095096          3    M       PortB     1     0      1
  2  9  111.537253          3    M       PortC     0     1      0
  3  20 104.713824          3    F       PortC     1     0      0
  4  28        NaN          2    F       PortA     1     1      0
```

```

▶ df.info()

→ <class 'pandas.core.frame.DataFrame'>
RangeIndex: 1000 entries, 0 to 999
Data columns (total 8 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   Age         1000 non-null    int64  
 1   Fare        960 non-null    float64 
 2   TravelClass 1000 non-null    int64  
 3   Gender       1000 non-null    object  
 4   EmbarkedPort 1000 non-null    object  
 5   SibSp        1000 non-null    int64  
 6   Parch        1000 non-null    int64  
 7   Survived     1000 non-null    int64  
dtypes: float64(1), int64(5), object(2)
memory usage: 62.6+ KB

▶ df.describe()

→          Age      Fare  TravelClass   SibSp   Parch  Survived
count  1.000000e+03  960.00000  1000.000000  1000.000000  1000.000000
mean   -4.611686e+17  93.505351   2.323000   1.21600   0.792000   0.349000
std    2.011193e+18  163.646527   0.788221   1.13958   0.918466   0.476892
min   -9.223372e+18  0.162020   1.000000   0.00000   0.000000   0.000000
25%   1.900000e+01   23.358816   2.000000   0.00000   0.000000   0.000000
50%   2.800000e+01   56.915283   3.000000   1.00000   1.000000   0.000000
75%   3.800000e+01   111.418688  3.000000   2.00000   1.000000   1.000000
max   1.100000e+02   2645.627902  3.000000   6.00000   5.000000   1.000000

```

```

▶ X = df.drop(['Survived'], axis=1)
y = df['Survived']

df.columns = df.columns.str.lower()
X.columns = X.columns.str.lower()
y.name = y.name.lower()

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f"Training data shape: {X_train.shape}")
print(f"Testing data shape: {X_test.shape}")

numeric_features_base = ['age', 'fare', 'sibsp', 'parch', 'travelclass']
categorical_features_base = ['gender', 'embarkedport']

preprocessor_base = ColumnTransformer(transformers=[
    ('num', SimpleImputer(strategy='median'), numeric_features_base),
    ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_features_base)
])

baseline_pipeline = Pipeline(steps=[
    ('preprocessor', preprocessor_base),
    ('classifier', LogisticRegression(random_state=42))
])

baseline_pipeline.fit(X_train, y_train)
y_pred_base = baseline_pipeline.predict(X_test)
baseline_accuracy = accuracy_score(y_test, y_pred_base)

print(f"\nBaseline Model Accuracy: {baseline_accuracy:.4f}")

→ Training data shape: (800, 7)
Testing data shape: (200, 7)

Baseline Model Accuracy: 0.3850

```

```

print("\n--- Part 2: Creating Features from Domain Knowledge ---")

X_train_eng = X_train.copy(deep=True)

X_train_eng['FamilySize'] = X_train_eng['sibsp'] + X_train_eng['parch'] + 1

X_train_eng['IsAlone'] = 0
X_train_eng.loc[X_train_eng['FamilySize'] == 1, 'IsAlone'] = 1

X_train_eng['FarePerPerson'] = X_train_eng['fare'] / X_train_eng['FamilySize']

X_train_eng['AgeGroup'] = pd.cut(
    X_train_eng['age'],
    bins=[0, 12, 18, 35, 60, 100],
    labels=['Child', 'Teen', 'Adult', 'MiddleAge', 'Senior']
)

print("Example of new features:")
print(X_train_eng[['age', 'fare', 'FamilySize', 'IsAlone', 'FarePerPerson', 'AgeGroup']].head())

print("\nUnique Age Groups found:")
print(X_train_eng['AgeGroup'].unique())

```

--- Part 2: Creating Features from Domain Knowledge ---

Example of new features:

age	fare	FamilySize	IsAlone	FarePerPerson	AgeGroup
179	7	74.847623	4	0	18.711906 Child
813	45	NaN	4	0	NaN MiddleAge
773	40	62.333398	1	1	62.333398 MiddleAge
428	63	48.132391	2	0	24.066196 Senior
592	37	23.567240	4	0	5.891810 MiddleAge

Unique Age Groups found:
 ['Child', 'MiddleAge', 'Senior', 'Adult', 'NaN', 'Teen']
 Categories (5, object): ['Child' < 'Teen' < 'Adult' < 'MiddleAge' < 'Senior']

▶ print("\n--- Part 3: Binning and Discretization ---")

```

X_train_binned = X_train.copy()
X_train_binned['age'].fillna(X_train_binned['age'].median(), inplace=True)

age_bins = [0, 12, 25, 60, 100]
age_labels = ['Child', 'Young Adult', 'Adult', 'Senior']
X_train_binned['AgeGroup'] = pd.cut(
    X_train_binned['age'],
    bins=age_bins,
    labels=age_labels,
    right=False
)

fare_bins = [0, 50, 100, 200, 500]
fare_labels = ['Low', 'Medium', 'High', 'Very High']
X_train_binned['FareGroup'] = pd.cut(
    X_train_binned['fare'],
    bins=fare_bins,
    labels=fare_labels,
    right=False
)

print("Example of 'AgeGroup' and 'FareGroup' features:")
print(X_train_binned[['age', 'AgeGroup', 'fare', 'FareGroup']].head())

print("\nValue counts for AgeGroup:")
print(X_train_binned['AgeGroup'].value_counts())

print("\nValue counts for FareGroup:")
print(X_train_binned['FareGroup'].value_counts())

```

```

→ --- Part 3: Binning and Discretization ---
Example of 'AgeGroup' and 'FareGroup' features:
   age AgeGroup      fare FareGroup
179    7   Child  74.847623   Medium
813   45   Adult     NaN       NaN
773   40   Adult  62.333398   Medium
428   63  Senior  48.132391     Low
592   37   Adult  23.567240     Low

Value counts for AgeGroup:
AgeGroup
Adult        464
Young Adult   213
Child         57
Senior        17
Name: count, dtype: int64

Value counts for FareGroup:
FareGroup
Low          341
Medium       210
High          153
Very High     57
Name: count, dtype: int64

```

```


▶ print("\n--- Part 4: Feature Transformation for Skewed Data ---")

sns.histplot(X_train['fare'], kde=True)
plt.title("Original Distribution of 'Fare'")
plt.xlabel("Fare")
plt.ylabel("Frequency")
plt.show()

fare_transformed = np.log1p(X_train['fare'].fillna(0))

sns.histplot(fare_transformed, kde=True, color='green')
plt.title("Log-Transformed Distribution of 'Fare'")
plt.xlabel("Log(Fare + 1)")
plt.ylabel("Frequency")
plt.show()

```



```

def create_custom_features(df):
    df_copy = df.copy()

    df_copy['FamilySize'] = df_copy['sibsp'] + df_copy['parch'] + 1
    df_copy['IsAlone'] = (df_copy['FamilySize'] == 1).astype(int)

    df_copy['FarePerPerson'] = df_copy['fare'] / df_copy['FamilySize']

    df_copy['AgeGroup'] = pd.cut(
        df_copy['age'],
        bins=[0, 12, 25, 60, 100],
        labels=['Child', 'Young Adult', 'Adult', 'Senior'],
        right=False
    )

    return df_copy

custom_feature_transformer = FunctionTransformer(create_custom_features, validate=False)

preprocessor_standard = ColumnTransformer(transformers=[

    ('num_impute_scale', Pipeline([
        ('imputer', SimpleImputer(strategy='median')),
        ('scaler', StandardScaler())
    ]), ['age', 'fare', 'sibsp', 'parch', 'FamilySize', 'IsAlone', 'FarePerPerson']),

    ('fare_log_impute_scale', Pipeline([
        ('imputer', SimpleImputer(strategy='median')),
        ('log', FunctionTransformer(np.log1p, validate=False)),
        ('scaler', StandardScaler())
    ]), ['fare']),

    ('cat_onehot', OneHotEncoder(handle_unknown='ignore'), ['gender', 'embarkedport', 'travelclass', 'AgeGroup'])
], remainder='drop')

full_pipeline = Pipeline(steps=[
    ('custom_features', custom_feature_transformer),
    ('preprocessor', preprocessor_standard),
    ('classifier', LogisticRegression(random_state=42, max_iter=1000))
])

full_pipeline.fit(X_train, y_train)
y_pred_full = full_pipeline.predict(X_test)
full_accuracy = accuracy_score(y_test, y_pred_full)

print(f"\nBaseline Model Accuracy: {baseline_accuracy:.4f}")
print(f"Full Feature-Engineered Model Accuracy: {full_accuracy:.4f}")

```

--- Part 5: Integrating Everything into a Full Pipeline ---
 Baseline Model Accuracy: 0.3850
 Full Feature-Engineered Model Accuracy: 0.6300

ANALYSIS AND DISCUSSION:

The dataset was successfully preprocessed using imputation, encoding, and transformation techniques. Automated workflows simplified data handling and improved model efficiency. Feature scaling and categorical encoding enhanced learning performance. Overall, preprocessing ensured cleaner data, reduced bias, and produced more accurate, reliable model predictions.